



Database Management Aggregate Functions

By: Meyer Tanuan (2022), Glenn Paulley (2015), John Mckay (2011)

Aggregate Functions

- Aggregate functions are used to count rows and summarize (typically) numeric data in a database
- Aggregate functions can appear in:
 - the column list of a SELECT statement
 - A HAVING clause (later in this unit)
 - An ORDER BY clause
- Aggregate functions differ from scalar functions
 - Scalar functions (e.g. ABS or SUBSTRING): one value in, one value returned
 - Aggregate functions: many values in, one value returned



The five basic aggregate functions

- Basic aggregate functions:
 - COUNT
 - AVG
 - SUM
 - MIN
 - MAX
- Each supports the ALL or DISTINCT keyword
 - ALL – compute the function over all non-NULL input values
 - AVG(ALL X)
 - DISTINCT – compute the result considering only unique non-NULL values
 - COUNT(DISTINCT X)

Aggregate Functions: COUNT

- COUNT always returns a value
 - COUNT(*) will return 0 (zero) if the input is empty
- All aggregate functions except COUNT(*) ignore NULL values in their input
 - For example, AVG(X) considers only non-null values of X
 - Aggregation functions other than COUNT will return NULL if they have all NULL values in their input



COUNT()

- COUNT(*) returns a count of the rows in a table
 - Result not related to nullability of any expression – merely a count of the rows in the input
- COUNT() always returns a value
 - If the input is empty, the result of COUNT(*) is 0 (zero)
- COUNT(expression)
 - Counts the number of non-null values in the input
- COUNT(DISTINCT expression)
 - Counts the number of distinct (duplicate-free) values

COUNT(*)

- COUNT(*) counts all rows unconditionally

```
SELECT COUNT( * ) AS [COUNT( * )]  
FROM Person
```

```
COUNT ( * )  
-----  
434
```

COUNT(expression)

- COUNT(expression) counts all rows with non-null values of expression

```
SELECT COUNT( fax ) AS [COUNT( fax )]  
FROM Employee
```

```
COUNT ( fax )
```

```
-----
```

```
4
```

COUNT(ALL)

- COUNT(ALL column) counts all rows with non-null values for column
 - ALL is the default and is usually omitted

```
SELECT COUNT( ALL fax ) AS [COUNT( ALL fax )]  
FROM Employee
```

```
COUNT ( ALL fax )
```

COUNT(DISTINCT) with WHERE

- Naturally, you can use a WHERE clause to restrict the rows being counted

```
SELECT COUNT( DISTINCT fax ) AS [COUNT( DISTINCT fax ) with WHERE]  
FROM Employee  
WHERE schoolCode = 'TAP' ;
```

```
COUNT ( DISTINCT fax ) with WHERE
```

AVG() – Compute the average

- AVG() is used to compute the average value of an expression

```
SELECT '$' + CONVERT(CHAR(7), CAST(AVG(amount) AS money),1)
      AS 'Average amount'
FROM Invoiceltem
```

Average amount

\$ 457.94

MIN() – Find the minimum value

- The MIN() aggregate function is used to find the minimum value from a set of values

```
SELECT '$' + CONVERT(CHAR(7), CAST(MIN(amount) AS money),1)
      AS [Minimum Price]
FROM Invoiceltem
```

Minimum Price

\$ 9.00

MAX() – Find the maximum value

- The MAX() aggregate function is used to find the maximum value from a set of values

```
SELECT '$' + CONVERT(CHAR(9), CAST(MAX(amount) AS money),1)
      AS [Maximum Price]
FROM Invoiceltem
```

Maximum Price

\$ 3,380.00

SUM() – Compute the sum

- The SUM(expression) aggregate function is used to compute the sum from all non-null values of expression

```
SELECT '$' + CONVERT(CHAR(9), CAST(SUM(amount) AS money),1)
      AS [Sum of all Items]
FROM InvoicItem
```

```
Sum of all Items
-----
$ 5,037.32
```

Aggregates with WHERE

- You can use a WHERE clause to restrict the input rows to the computation of AVG(), MIN(), MAX(), and SUM() as well

```
SELECT '$' + CONVERT(CHAR(10), CAST(AVG(amount) AS money),1)
    AS [Avg of all items]
FROM Invoiceltem
WHERE item LIKE '%CSI%'
```

Avg of all items

\$ 210.63

GROUP BY and HAVING

GROUP BY

- GROUP BY, in combination with aggregate functions, can be used to group rows and summarize the results of aggregate functions for each group
- Once GROUP BY is used, the expressions in a query's SELECT list are restricted to:
 - Expressions specified in the GROUP BY clause
 - Aggregate functions (or other computations based on their results)



What does GROUP BY do?

- GROUP BY partitions the input rows (FROM and WHERE clauses) into groups that have the identical set of values for the GROUP BY expressions specified in the query
- Once all of the input rows are partitioned, the server then computes the result of the aggregate functions for each partition
- Each partition then generates a single row in the output



Using GROUP BY

```
SELECT studentNumber  
       , SUM(amount)  
FROM Payment  
GROUP BY studentNumber
```

| studentNumber | |
|---------------|----------|
| ----- | ----- |
| 1114453 | 1000.00 |
| 1335314 | 10000.00 |
| 2286425 | 1000.00 |
| 2642726 | 13159.00 |
| 2826147 | 1000.00 |
| 3244449 | 5291.44 |
| 3868247 | 13159.00 |
| 6644710 | 1159.00 |
| 7252620 | 5291.44 |
| 7677479 | 2000.00 |
| 7681752 | 1000.00 |
| 7826662 | 6791.44 |
| 8431710 | 5291.44 |
| 8588766 | 1000.00 |
| 8866782 | 3000.00 |

What forms a GROUP?

- Groups are formed from unique values of the combination of grouping columns in the GROUP BY clause
 - For the purposes of comparing column or expression values, NULL is semantically equivalent to NULL, as in SELECT DISTINCT
- GROUP BY X
 - Groups formed based on unique values of X
- GROUP BY X, Y
 - Groups formed based on unique combinations of values of X and Y for the same row

Empty input with GROUP BY

- If a query's intermediate result is empty, and that query contains GROUP BY, then the query also returns the *empty set* (i.e., 0 rows affected)

```
SELECT sessionCode, AVG(amount) AS AverageAmount
FROM Invoice
WHERE amount > 0
GROUP BY sessionCode
```

```
sessionCode  AverageAmount
```

```
-----
```

```
(0 rows affected)
```

Empty input without GROUP BY

- For queries that involve aggregation but don't have a GROUP BY clause, then the query does not return an empty set (i.e., 1 row affected)

```
SELECT AVG(amount) AS AverageAmount  
FROM Invoice  
WHERE amount > 0
```

```
AverageAmount
```

```
-----
```

```
NULL
```

```
(1 row affected)
```

HAVING

- HAVING can be used to limit the results that appear in groups created with GROUP BY
- HAVING is similar to a WHERE clause
 - WHERE restricts the rows created by the FROM clause
 - HAVING restricts the groups created from the GROUP BY clause
- HAVING is often used with aggregate functions
 - You cannot use an aggregate function in a WHERE clause because at the point of applying the WHERE conditions, the groups have yet to be formed



GROUP BY and HAVING

- Determine the SUM of payments for specific students but only for students who have paid over more than one payment

```
SELECT studentNumber  
      , '$' + CONVERT( CHAR(12), CAST( SUM(amount) AS money ), 1 )  
      AS [Invoice Total]  
FROM Payment  
GROUP BY studentNumber  
HAVING COUNT(*) > 1
```

GROUP BY and HAVING Result

| studentNumber | Invoice Total |
|---------------|---------------|
| 2642726 | \$ 13,159.00 |
| 7677479 | \$ 2,000.00 |
| 7826662 | \$ 6,791.44 |
| 8431710 | \$ 5,291.44 |
| 8866782 | \$ 3,000.00 |

ORDER OF CLAUSES

- If a query contains a WHERE clause, you must put the WHERE clause before GROUP BY and HAVING
- ORDER BY follows GROUP BY and the optional HAVING clause

```
SELECT studentNumber
      , '$' + CONVERT( CHAR(12), CAST( SUM(amount) AS money ), 1 )
      AS [Invoice Total]
FROM Payment
WHERE studentNumber <> 8431710
GROUP BY studentNumber
HAVING COUNT(*) > 1
ORDER BY 2 DESC;
```

GROUP BY, HAVING, WHERE and ORDER BY Result

| studentNumber | Invoice Total |
|---------------|---------------|
| 2642726 | \$ 13,159.00 |
| 7826662 | \$ 6,791.44 |
| 8866782 | \$ 3,000.00 |
| 7677479 | \$ 2,000.00 |

Limitation: Expressions with GROUP BY

- One cannot specify an expression in a query's SELECT list that is NOT an aggregate function unless the identical expression appears in the query's GROUP BY clause
- For example:
 - SELECT X, SUM(Y) FROM T GROUP BY X – permitted
 - SELECT X, SUM(Y) FROM T – illegal
 - SELECT X, Y, SUM(Z) FROM T GROUP BY X – illegal
 - SELECT X, SUM(Z) FROM T GROUP BY X, Y – permitted



Example: GROUP BY WITH ROLLUP

- The ROLLUP operator in a query adds a final summary row to the result (SQL Server 2008 or later)
- Select all persons grouping them by first letter of their last name and provide a count of the number of persons with the same letter of their last name (e.g., 24 unique first letters)

```
SELECT LEFT(LastName,1), COUNT(LEFT(LastName,1))  
FROM Person  
GROUP BY LEFT(LastName,1) WITH ROLLUP
```

Example: Add search condition (7 letters dropped)

- Select all international students grouping them by first letter of their last name and provide a count of the number of students with each first letter (e.g., 17 unique first letters)

```
SELECT LEFT(LastName,1), COUNT(LEFT(LastName,1))  
FROM Student s JOIN Person p  
    ON s.number = p.number  
WHERE isInternational=1  
GROUP BY (LEFT(LastName,1)) WITH ROLLUP
```

Example: OUTER JOIN and join table expression

- To show all the original 24 unique first letters (some with zero)
 - RIGHT OUTER JOIN
 - Use join table expression (move search condition from WHERE clause to ON clause)

```
SELECT LEFT(LastName,1), COUNT(s.number)
FROM Student s RIGHT OUTER JOIN Person p
    ON (s.number = p.number AND isInternational=1)
GROUP BY (LEFT(LastName,1)) WITH ROLLUP
```