



# Database Management Create Tables

By: Meyer Tanuan (2022), Rick Kozak (2019), Glenn Paulley (2015), John Mckay (2011)

# SQL Data Types

- Each column in a table has a data type
- The data type determines the type and amount of data you can store in a column
- Different types have different semantics and support different sets of operations

# Character Types

## ASCII

- CHAR(n) – Fixed length max 8K
- VARCHAR(n) – Variable length max 8K
- TEXT – Variable length max 2GB

## Unicode

- NCHAR(n) – Fixed length max 4000
- NVARCHAR(n) – Variable length max 2GB
- NTEXT – Variable length max  $2^{30}-1$

# Binary Types

- BINARY(n) – fixed length max 8K bytes
- VARBINARY(n) – variable length max 8K
- IMAGE – variable length max 2GB
- BIT – 0, 1, or NULL

# Exact Numeric Types

- INT – 32 bit signed number
- BIGINT – 64 bit signed number
- SMALLINT – 16 bit signed number
- TINYINT – 8 bit unsigned number

# Decimal Numeric Types

- DECIMAL and NUMERIC are synonyms
- DECIMAL is an exact numeric type that supports numeric values with fractional component
- Declaration specifies *precision* (between 1 and 38) and an optional *scale*
- DECIMAL(7,2)
  - 7 significant digits (precision), 2 of these digits after the decimal point (scale)
- DECIMAL(10)
  - 10 significant digits, no fractional component
- Range:  $\pm 10^{38}$
- Storage: between 5 and 17 bytes
- Default: if a value is declared DECIMAL, assumed to be DECIMAL(18)

# Approximate Numeric Types

- Approximate numeric types store floating-point values; specification uses scientific notation (mantissa and exponent)
- Not all numbers can be represented exactly (mantissa is in base-2)
- Should never be used for primary keys
- REAL
  - Storage: 4 bytes
  - Range:  $-3.40E + 38$  to  $-1.18E -38$ , 0 and  $1.18E -38$  to  $3.40E + 38$
- FLOAT(n)
  - Value of n is between 1 and 53; number of bits used to store the mantissa of the number
  - Storage: 4 bytes if  $n \leq 24$ ; 8 bytes if  $n > 24$
- DOUBLE PRECISION
  - Synonym for FLOAT(53)

# Date Type

- Range: 1 Jan 0001 AD - 31 Dec 9999
  - Cannot be used for dates BC
- Date only, no time portion
- Default string format: 'YYYY-MM-DD'
- Storage size: 3 bytes
- Default value: 1 January 1900



# Time Type

- 00:00:00.00000000 - 23:59:59.99999999
- A TIME type can have msec precision
- Default is TIME(7), 100ns precision
- Time only, no date portion
- Default string format: 'hh:mm:ss[.nnnnnnnn]'
- Storage size: between 3 and 5 bytes, depending on precision
- Default value: 00:00:00

# DateTime Type

- Date range: Jan 1, 1753 - Dec 31, 9999
- Time range: 00:00:00 and 23:59:59:997
  - Approximately 1/300 second precision
- Character length: 19 to 23 characters
- String: 'YYYY-MM-DD hh:mm:ss[.nnnnnnnn]'
- Storage size: 8 bytes
- Default value: 1900-01-01 00:00:00.000
- SET DATEFORMAT option controls interpretation of DATETIME values when specified as string literals

# SmallDateTime Type

- Date range: January 1, 1900 - June 6, 2079
- Time range: 00:00:00 - 23:59:59
  - Approximately one minute precision
- Character length: 19 characters maximum
- Default format: 'YYYY-MM-DD hh:mm:ss'
- Storage size: 4 bytes
- Default value: 1900-01-01 00:00:00

# DateTime2 Type

- Equivalent to the TIMESTAMP type from the ISO SQL Standard
- Now recommended for all applications instead of DATETIME
- Date range: January 1, 0001 - December 31, 9999
  - Cannot be used for dates BC
- Time range: 00:00:00 - 23:59:59:99999999
- DATETIME2 allows time precision, up to 100ns
  - DATETIME2(7) is the default (100ns precision)
- Character length: 19 to 27 characters
- Default string format: 'YYYY-MM-DD hh:mm:ss[.nn]'
- Storage size: 6 - 8 bytes depending on time precision
- Default value: 1900-01-01 00:00:00.00000000

# DateTimeOffset Type

- Equivalent to the TIMESTAMP WITH TIME ZONE type from the ISO SQL Standard
- Adds a time zone component to a DATETIME2 type
- Date range: January 1, 0001 - December 31, 9999
  - Cannot be used for dates BC
- Time range: 00:00:00 - 23:59:59:99999999
  - allows time precision, up to 100ns
- DATETIMEOFFSET(7) is the default (100ns precision)
  - Time zones: -14h through +14h
  - Character length: 26 - 34 characters
  - Format: YYYY-MM-DD hh:mm:ss.nnnnnnnn{+|-}hh:mm
  - Storage size: 10 bytes
- Default value: 1900-01-01 00:00:00.00000000 +00.00

# Other Types

- **MONEY**

Money amounts to the trillions with 4 decimal places accuracy

- **ROWVERSION**

- A database-wide unique row number

# CREATE TABLE

- Basic CREATE TABLE syntax and format:

```
CREATE TABLE TableName  
(column1 TYPE [qualifiers],  
  column2 TYPE,  
  column3 TYPE [qualifiers],  
  . . .  
  [, table-level constraints]  
)
```

- It is customary, but not essential, to put primary key columns at the beginning of a table

## **CREATE TABLE: Example**

```
CREATE TABLE Product
(
    productId VARCHAR(11),
    description VARCHAR(75),
    vendorId VARCHAR(4),
    vendorPartNumber VARCHAR(20),
    price DECIMAL(7,2),
    reorderThreshold INT,
    productCategoryCode CHAR(1)
)
```



# NULL , NOT NULL

Each column definition can include:

- NULL to indicate that NULL values are permitted
  - This is the default and is therefore seldom seen
- NOT NULL to indicate that NULL values are prohibited

`vendorName VARCHAR(50) NOT NULL,`

# CREATE TABLE

All database systems have various optional clauses on CREATE TABLE that make these statements non-portable across various database management systems

- CREATE TABLE can contain vendor-specific syntax for:
  - Column and/or table-level encryption
  - Column compression
  - FOREIGN KEY constraints and various options
  - Table partitioning specification
  - Indexing options
  - Use of system-generated values for primary keys
  - SQL Server supports both IDENTITY values and SEQUENCE types

# SELECT INTO

```
SELECT <select list>  
INTO <table>  
FROM ...
```

- The columns in <table> have the same names and data types as the attributes in the query's SELECT list
- Rows are populated from the result of the SELECT statement

# **SELECT INTO Limitations**

- Cannot be used to create a partitioned table
- Indexes, constraints, and triggers defined in the source table are not transferred to the new table, nor can they be specified in the SELECT...INTO statement.
  - If these objects are required, you can create them after executing the SELECT...INTO statement.
- Specifying an ORDER BY clause does not guarantee the rows are inserted in the specified order.
- When a computed column is included in the select list, the corresponding column in the new table is not a computed column. The values in the new column are the values that were computed at the time SELECT...INTO was executed.

# DROP TABLE

DROP TABLE <table>

- Completely removes table from database

## IF EXISTS

Use code like this in a script to drop a table if it already exists:

```
IF EXISTS (  
    SELECT name  
    FROM sysobjects  
    WHERE name = 'Mytable')  
DROP TABLE Mytable;
```

# CONSTRAINTS

- A. Primary Key (PK)
- B. Unique (UX)
- C. Foreign Key (FK)
- D. Check (CK)
- E. Default (DF)

# CONSTRAINTS

## Types

### DEFAULT

- The value stored in a column on INSERT if no value is specified

### PRIMARY KEY

- Identifies one or more columns that uniquely identify each row in the table
- All values in a PRIMARY KEY cannot be NULL

### UNIQUE

- Prohibits duplicate values

### FOREIGN KEY

- Used to enforce referential integrity

### CHECK

- Used to enforce domain integrity



# CONSTRAINTS: Naming

- Constraints should be named
- If you don't name a constraint, the database engine synthesizes a long and ugly name
- Often, by convention, PRIMARY KEY, UNIQUE KEY, FOREIGN KEY, DEFAULT and CHECK constraint names **start** or **end** in PK, UX, FK, DF and CK respectively, for example:

```
CONSTRAINT CK_Guests_NameLength  
CHECK ...
```

# CONSTRAINTS

## Categories

Constraints are one of two categories:

- Column constraint
  - A constraint that applies to one and only one column. Often NOT NULL is specified as a column constraint for any given column
- Table constraint
  - A constraint specified at the table level
  - Must be used if the constraint involves multiple columns, for example a multi-column PRIMARY KEY or FOREIGN KEY constraint

## TABLE CONSTRAINT: Example

All constraints *can* be specified at the table level

- Multi-column constraints *must* be done at the table level
- Table constraint example:

```
CREATE TABLE OrderItem
(
    invoiceNumber INT NOT NULL,
    productId VARCHAR(11) NOT NULL,
    quantity INT,
    discount NUMERIC(3,3)
    CONSTRAINT PK_OrderItem
    PRIMARY KEY (invoiceNumber,
                 productId)
)
```

# CONSTRAINTS: Examples

```
CREATE TABLE Product
(
    productId VARCHAR(11) NOT NULL
    CONSTRAINT PK_Product PRIMARY KEY,

    price NUMERIC(7,2)
    CONSTRAINT DF_Product_Price DEFAULT 0.00,

    reorderThreshold INT
    CONSTRAINT CK_Product_Threshold_GtZero
    CHECK (reorderThreshold > 0),

    categoryCode CHAR(1)
    CONSTRAINT CK_Product_CategoryCode_valid
    CHECK (categoryCode IN ('S','H'))
)
```

# **CONSTRAINTS:**

## **A. PRIMARY KEY**

- PRIMARY KEY values cannot be NULL
  - Typically one ensures this by adding NOT NULL constraints to each primary key column
- A primary key uniquely identifies each row in a table

# CONSTRAINTS:

## B. UNIQUE

A UNIQUE constraint differs slightly from a PRIMARY KEY:

- A table can have only one PRIMARY KEY, but many UNIQUE constraints
- UNIQUE constraints can specify nullable columns
  - Only one (1) NULL value is permitted in the index
  - This restriction is not enforced by other DBMS products, for example SAP SQL Anywhere

As with PRIMARY KEYS:

- A UNIQUE constraint specification results in the creation of a unique index on those columns
- UNIQUE constraints can be referenced by FOREIGN KEY constraints in other tables

## **CONSTRAINTS:**

### **C. FOREIGN KEY**

- A FOREIGN KEY constraint provides referential integrity for the data in the column(s).
- FOREIGN KEY constraints require that each value in the column exist in the specified column(s) in the referenced table.
- The referenced columns in the referenced table are typically those specified in that table's PRIMARY KEY constraint

# Referential Integrity

- Referential integrity (or RI) means that a row cannot be inserted into the referencing table unless a row with matching values exists in the referenced table
  - Often this is referred to as a “parent” and “child” relationship amongst tables in a schema

Example:

```
INSERT INTO Parent VALUES(1, 'Fred'); -- OK
INSERT INTO Child VALUES( 1 ); -- OK
INSERT INTO Child VALUES( 2 ); -- Error
```



## Foreign Key Constraint Options

SQL Server supports the following options for ON DELETE and ON UPDATE rules:

- **NO ACTION:** the operation is not permitted. In other systems this is called RESTRICT, as in “ON DELETE RESTRICT”. This is the default.
- **CASCADE:** the (delete or update) operation is “cascaded” to the referencing table. Rarely used in practice.
- **SET NULL:** Rather than leave orphaned rows in the child table, the column value(s) in the referencing row are set to NULL.
- **SET DEFAULT:** Similar to SET NULL, but instead the default values for the columns in the child table are used.

# CONSTRAINTS:

## D. CHECK

- A CHECK constraint is a constraint that the database verifies with each INSERT, UPDATE, or MERGE statement that affects the columns referenced by the constraint
- CHECK constraints are not verified by DELETE operations (why?)
- CHECK constraints are violated only when they evaluate to FALSE
- A CHECK constraint that evaluates to UNKNOWN is considered valid

# CONSTRAINTS: CHECK Example

Example:

`CHECK( x < 15)`

- In this case, the server will verify that the X-value in each row is less than 15 after each update modification
- Otherwise the INSERT, UPDATE, or MERGE statement will get an error
- Expressions in CHECK constraints (X in this example) refer to column values in the particular row being modified
- However: CHECK constraints support sub-queries – hence the search condition in the CHECK constraint may be as complex as desired

## CONSTRAINTS:

### E. DEFAULT

In practice, the DEFAULT constraint is often not named.

- You may see either of these alternatives:

```
price NUMERIC(7,2)  
CONSTRAINT DF_Product_Price  
DEFAULT 0.00
```

```
price NUMERIC(7,2)  
DEFAULT 0.00
```

- DEFAULT constraints are useful for situations where not all values are known at the time a row is inserted, but is desirable that the column not contain NULL

# Requirements for Keys (1)

## Uniqueness

- An obvious requirement for a key
- Names, phone numbers are usually inappropriate, though still serve as useful, if not essential, search terms
  - Bell Canada: maintains both a client ID, and an account number, along with phone numbers

## Control

- Is the key value under the application's complete control? Or is the value under the control of an independent organization, for example the Province of Ontario or the Government of Canada?

# Requirements for Keys (2)

## Immutability

- Updating primary keys is rarely a good idea
  - System maintenance can be complex – need to update all related tuples in other tables
  - Potential loss of historical information

## Data type

- Alphabetic keys have implications for internationalization
- Never use floating point values for keys

# Requirements for Keys (3)

## Entropy and range

- Running out of “room” is problematic because it necessitates both schema and application changes
- How many bits does it take to store the key?
- How many bits are required to effect a change in the key value?
- How many bits are different between the letters ‘a’ and ‘b’ in 7-bit ASCII?

## Generation technique

- System-generated or application-generated?
- There are significant performance implications given the choices because an application-generated key will likely require additional SQL requests to insert a new row

# Requirements for Keys (4)

## Format

- Different key formats can be useful to differentiate between business entities
  - GHK009 – client identifier (note: alphanumeric; Latin characters; no vowels)
  - 897765 – group insurance policy number

- Fast, easy differentiation is useful

## Self-checking

- Self-checking keys can aid in administering client data
- Examples: American Express numbers, Canadian SIN numbers
- SINs in Canada: 8 digits plus a “check” digit
  - Computed using Luhn’s algorithm: see [http://en.wikipedia.org/wiki/Luhn\\_Algorithm](http://en.wikipedia.org/wiki/Luhn_Algorithm)



# Requirements for Keys (5)

## Embedded information

- Often embedded information is held within an identifier; SIN is one example
  - First digit of a Canadian SIN number indicates the region in which the number was issued

## Input difficulty

- Can help to reduce data entry errors; example: Canadian postal codes
  - Very difficult to touch-type a postal code even for superb typists
  - Significant reduction in misdirected mail

# Requirements for Keys (6)

## Composite or simple

- Simple keys of a single column are significantly easier to deal with in an application
- Simplifies the logical and physical schema
  - uses less storage

# IDENTITY

- IDENTITY columns are typically used as primary keys
- The auto-generated value becomes a surrogate key for the row
- SQL Server generates a unique, sequential, integer identity value for each row inserted in a table
- IDENTITY columns are usually INTEGER but they can be stored using other numeric data types

## IDENTITY: Example

```
CREATE TABLE Wedding
(
    id INT IDENTITY NOT NULL
    CONSTRAINT PK_Wedding
    PRIMARY KEY,
    partner1 VARCHAR(50) NOT NULL,
    partner2 VARCHAR(50) NOT NULL,
    date SMALLDATETIME NOT NULL,
    location VARCHAR(30) NOT NULL
)
```

# INSERT with IDENTITY

- Usually, you must omit the identity column from the INSERT statement
- This is because SQL Server is responsible for generating key values, not you
- However, you can explicitly enable the ability to override this behaviour and INSERT explicit values to IDENTITY columns
- This is done using the SET IDENTITY\_INSERT statement:

```
SET IDENTITY_INSERT Guest ON;
```

# INSERT with IDENTITY: Examples

```
INSERT INTO wedding  
(partner1, partner2, date, location)  
VALUES  
('Jane', 'Bob', '2022-06-01', 'Elmira')
```

```
INSERT INTO wedding  
(partner1, partner2, date, location)  
VALUES  
('Eli', 'Betsy', '2021-12-15', 'St.  
Jacobs')
```

```
INSERT INTO wedding  
(partner1, partner2, date, location)  
VALUES  
('Pat', 'Bruce', '2021-06-30', 'Toronto')
```

# INSERT with IDENTITY: Result

```
SELECT *  
FROM wedding;
```

id	partner1	partner2	date	location
1	Jane	Bob	2022-06-01	Elmira
2	Eli	Betsy	2021-12-15	St. Jacobs
3	Pat	Bruce	2021-06-30	Toronto

# Last IDENTITY

The most common way to find the last identity value inserted is with @@IDENTITY:

```
SELECT @@IDENTITY
```

You may want to use an alias:

```
SELECT @@IDENTITY  
AS 'Last Identity Value'
```

Last Identity Value

-----

3



# Last IDENTITY: Three Ways

There are actually three ways to find the last identity value:

## `@@IDENTITY`

- Returns the identity value generated by the previous INSERT statement

## `SCOPE_IDENTITY( )`

- Same as `@@IDENTITY`, but restricted to the current scope (useful for stored procedures)

## `IDENT_CURRENT( table )`

- Returns the last identity value created by any INSERT in a particular table

# Last IDENTITY in SELECT

You can use @@IDENTITY in a WHERE clause to work with the last row inserted

```
SELECT *  
FROM Wedding  
WHERE id = @@IDENTITY
```

# IDENTITY: IDENT\_CURRENT Example

```
CREATE TABLE Guest (  
  id INT IDENTITY PRIMARY KEY,  
  name NVARCHAR(10) NOT NULL)
```

```
CREATE TABLE Reservation (  
  id INT IDENTITY PRIMARY KEY,  
  guestId INT CONSTRAINT  
  FK_Reservation_Guest FOREIGN KEY  
  REFERENCES Guest(id),  
  checkInDate DATE NOT NULL)
```

```
DECLARE @GuestId INT;  
INSERT INTO Guest VALUES('Meyer');  
SET @GuestId =  
  IDENT_CURRENT('Guest');  
INSERT INTO Reservation VALUES  
  (@GuestId, '2022-01-01');
```

## IDENTITY Variation

- Identity column may start with a different value than 0 and increment with a value other than 1, but this is seldom done in practice

```
<column> INTEGER  
NOT NULL IDENTITY(15,2)
```

- Initial value of 15, increment by 2

# UNIQUE IDENTIFIER

SQL Server UNIQUEIDENTIFIER columns are used to store Globally Unique Identifier (GUID) values

- GUIDs are made up of 32 Hexadecimal (base 16) numbers, for example:

cd2beaf3-2203-483d-b5a6-8eac46cf8c16

- GUIDs occupy 128 bits / 16 bytes

# UNIQUE IDENTIFIER (2)

- The probability of 2 UNIQUEIDENTIFIERs having the same value is miniscule
- Many applications treat UNIQUEIDENTIFIER values as unique, without fear of collisions
- UNIQUEIDENTIFIER can be used for the primary key to avoid collisions when data from different sources are combined
- Issues
  - GUID fields are large (16 bytes) compared to INTEGER values
  - Not self-checking
  - Difficult to type, and easy to transpose adjacent characters

# UNIQUE IDENTIFIER (3)

- SQL Server has two functions for creating UNIQUEIDENTIFIER values:

**NEWID()**

- The original function

**NEWSEQUENTIALID()**

- A function introduced in SQL Server 2014 that makes sure each new value is larger than any previous value on a given computer until the next restart of the Windows operating system

# UNIQUE IDENTIFIER: NEWSEQUENTIALID Example

```
CREATE TABLE Software
(
    uniqueId UNIQUEIDENTIFIER NOT NULL

    CONSTRAINT DF_Software_UniqueId
    DEFAULT NEWSEQUENTIALID()

    CONSTRAINT PK_Software
    PRIMARY KEY (uniqueId),

    product NVARCHAR(50) NOT NULL,

    ...
);
```



# SEQUENCE

- A SEQUENCE is a database object that generates a sequence of numeric values and are, like IDENTITY columns, used to create surrogate key values
- Advantages:
  - More flexible than IDENTITY
  - Cycling of values can be controlled
  - The sequence value is known **prior** to the execution of the INSERT statement
  - Sequence values can be generated in descending order
  - Can apply to more than one table; a SEQUENCE is an object in the database separate from any table

# CREATE SEQUENCE Syntax

```
CREATE SEQUENCE [schema_name.]sequence_name  
[AS[integer_type]]  
[START WITH <constant>]  
[INCREMENT BY <constant>]  
[{MINVALUE [<constant>]}|{NOMINVALUE}]  
[{MAXVALUE [<constant>]}|{NOMAXVALUE}]  
[CYCLE|{NOCYCLE}]  
[{CACHE [<constant>]}|{NO CACHE}]
```

# SEQUENCE: Integer Types

- A sequence can be defined as any integer type:
  - **Tinyint** - Range 0 to 255
  - **Smallint** - Range -32,768 to 32,767
  - **Integer** - Range -2,147,483,648 to 2,147,483,647
  - **Bigint** - Range -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
- Decimal and numeric with a scale of 0.
- Any user-defined data type (alias type) that is based on one of the allowed types.
- If no data type is provided, the **bigint** data type is used as the default.

## SEQUENCE: Example

```
CREATE SEQUENCE CountBy1  
START WITH 1  
INCREMENT BY 1;
```

```
CREATE SEQUENCE CountByNeg1  
START WITH 0  
INCREMENT BY -1;
```

# SEQUENCE: NEXT VALUE FOR

- To retrieve the next value from a SEQUENCE object, you code a SELECT statement using the NEXT VALUE FOR function

```
SELECT NEXT VALUE FOR  
Test.CountBy1
```

- NEXT VALUE FOR is a function that can be used in any SELECT statement
- Note that NEXT VALUE FOR has many restrictions and cannot be used in a query's WHERE clause
- If there are multiple NEXT VALUE FOR functions over the same sequence in the same SELECT statement, all of the NEXT VALUE FOR functions will return the same value

# SEQUENCE: NEXT VALUE FOR Example

A typical use case for a SEQUENCE is to generate surrogate keys for newly-inserted rows:

```
CREATE TABLE TestTable
(
    counterColumn INT PRIMARY KEY,
    name VARCHAR(40) NOT NULL
);

INSERT INTO TestTable
( counterColumn, name )
VALUES
( NEXT VALUE FOR CountBy1, 'x' );
```

# Database Management

## Alter Tables

By: Meyer Tanuan (2022)

*Source: Murach's SQL Server 2019 for Developers*

# Basic Syntax of the **ALTER TABLE** statement

```
ALTER TABLE table_name [WITH CHECK|WITH NOCHECK]
{ADD new_column_name data_type
[column_attributes] |
DROP COLUMN column_name |
ALTER COLUMN column_name new_data_type
[NULL|NOT NULL] |
ADD [CONSTRAINT] new_constraint_definition |
DROP [CONSTRAINT] constraint_name}
```



# ALTER TABLE Examples (1)

-- Add a column

```
ALTER TABLE Vendors  
ADD LastTranDate DATE NULL;
```

-- Drop a column

```
ALTER TABLE Vendors  
DROP COLUMN LastTranDate;
```

-- Add a new check constraint

```
ALTER TABLE Invoices WITH NOCHECK  
ADD CHECK (InvoiceTotal >= 1);
```

## ALTER TABLE Examples (2)

-- Add a foreign key constraint

```
ALTER TABLE InvoiceLineItems WITH CHECK  
ADD FOREIGN KEY (AccountNo)  
REFERENCES GLAccounts(AccountNo);
```

-- Change the data type of a column

```
ALTER TABLE InvoiceLineItems  
ALTER COLUMN InvoiceLineItemDescription  
VARCHAR(200);
```